



Image

HOW TO CREATE A CHESS ENGINE USING DEEP REINFORCEMENT LEARNING

A CRITICAL LOOK AT DEEPMIND'S ALPHAZERO

INTERNAL PROMOTOR: WOUTER GEVAERT

EXTERNAL PROMOTOR: <NAME HERE>

RESEARCH CONDUCTED BY

TUUR VANHOUTTE

FOR OBTAINING A BACHELOR'S DEGREE IN

MULTIMEDIA & CREATIVE TECHNOLOGIES

HOWEST | 2021-2022

Preface

This bachelor thesis is the conclusion to the bachelor program Multimedia & Creative Technologies at Howest college West Flanders in Kortrijk, Belgium. The program teaches students a wide range of skills in the field of computer science, with a focus on creativity and Internet of Things. From the second year on, students can choose between four different modules:

1. **AI Engineer**
2. **Smart XR Developer**
3. **Next Web Developer**
4. **IoT Infrastructure Engineer**

This bachelor thesis was made under the **AI Engineer** module. The subject of the thesis is a critical look at the result of my research project in the previous semester. The goal of the project was to create a chess engine in Python with deep reinforcement learning based on DeepMind's AlphaZero algorithm.

I will explain the research I needed to create it, the technical details on how to program the chess engine and I will reflect on the results of the project. To do this, I will contact multiple people and communities familiar with the field of reinforcement learning to get a better understanding of the impact of this research on society. Based on this, I will give advice to people and companies who wish to implement similar algorithms.

I would like to show gratitude to Wouter Gevaert for his enthusiastic support in the creation of my research project and this thesis. I also want to thank the other teachers at Howest Kortrijk, who shared their knowledge and expertise in programming and AI in very interesting classes.

Furthermore, I would like to thank my parents for giving me the chance to have a good education, and the motivation to get the best I can out of my studies.

Tuur Vanhoutte, 1st June 2022

Abstract

This bachelor thesis answers the question: “How to create a chess engine using deep reinforcement learning?”. It explains the difference between normal chess engines and chess engines that use deep reinforcement learning, and specifically tries to recreate the results of AlphaZero, the chess engine by DeepMind, in Python on consumer hardware.

The technical research shows what is needed to create my implementation using Python and TensorFlow. It shows how to program the chess engine, how to build the neural network, and how to train and evaluate the network. During the creation of this chess engine, it was crucial to create a huge amount of data through self-play.

The thesis contains a reflection on the results of my research project, which proposes a solution to the problem of creating a high amount of games through self-play. It also reflects on the impact of this research on society, and the viability of this type of artificial intelligence in the future. It follows with advice to people and companies who wish to implement similar algorithms.

Contents

Preface	1
Abstract	3
Contents	5
List of figures	7
List of tables	8
List of abbreviations	8
Glossary	9
1 Introduction	10
2 Research	11
2.1 What is a chess engine?	11
2.2 How do traditional chess engines work?	11
2.2.1 The minimax algorithm	11
2.2.2 The evaluation function	11
2.2.3 Pseudocode	12
2.2.4 Alpha-beta pruning	12
2.3 Monte Carlo Tree Search	12
2.3.1 The algorithm	13
2.4 Go	13
2.4.1 AlphaGo	14
2.4.2 AlphaGo Zero	14
2.5 AlphaZero	14
2.5.1 Neural network input	15
2.5.2 Neural network layers	15
2.5.3 Neural network output	16
2.6 Training the network	16
2.6.1 Tensor Processing Units (TPU)	16
2.7 Leela Chess Zero	17
3 Technical research	18
3.1 Introduction	18
3.2 Class structure	18
3.2.1 Making one move	19
3.3 The neural network	20
3.4 A tree structure with nodes and edges	20
3.5 The MCTS algorithm	21
3.5.1 The selection step	21
3.5.2 The expansion step	22
3.5.3 The evaluation step	23
3.5.4 The backpropagation step	23
3.6 Creating the dataset	24
3.6.1 Creating a dataset from puzzles	24
3.7 Training the neural network	25

3.7.1	The first training session	26
3.7.2	The second training session	26
3.8	Multiprocessing	27
3.8.1	Without multiprocessing	27
3.8.2	With multiprocessing	28
3.9	A GUI to play against the engine	28
3.10	Docker images	30
3.11	The final project	30
3.11.1	Creating your own untrained AI model	30
3.11.2	Creating a training set through self-play	31
3.11.3	Evaluating two models	31
3.11.4	Playing against the AI	32
3.12	Porting to C++	32
4	Reflection	33
4.1	Introduction	33
4.2	Strengths and weaknesses	33
4.3	Is the result of the project usable in the corporate world?	34
4.4	What are the possible obstacles for companies that wish to implement this?	34
4.5	What is the added value for companies?	34
4.6	Which alternatives are there?	34
4.7	Is there a socio-economic impact present?	35
4.8	Is there opportunity for further research?	35
5	Advice	36
5.1	Introduction	36
5.2	When should you use this technology?	36
5.3	Recommendations	36
5.4	Tips when programming a similar application	36
5.4.1	Creating the tree data structure	36
5.4.2	Programming the MCTS algorithm	36
5.4.3	Choosing the right input for the neural network	37
6	Conclusion	38
7	Bibliography	39
8	Appendix	42

List of figures

1	Depth-First search vs Breadth-First search [6]	11
2	The 4 steps of the MCTS algorithm [10]	13
3	Go board [12]	14
4	Basic class structure for the code responsible for playing a game	18
5	Pipeline to make one move	19
6	Example tree after 400 simulations. N = amount of times the selection step selects that edge	20
7	The four steps in AlphaZero's MCTS algorithm [29]	21
8	Board example	22
9	The input state converted from the above example board	22
10	A subset of the 73 8x8 planes from the policy output. The brighter the pixel, the better the move.	23
11	The same subset, but with the illegal moves filtered out	23
12	Some output planes from an untrained model.	23
13	The training set consists of: the input state, the move probabilities, and the eventual winner .	24
14	The training pipeline	25
15	First training session	26
16	Second training session	26
17	Self-play without multiprocessing	27
18	Self-play with multiprocessing	28
19	Chessboard example when playing against the engine	29
20	Promoting a pawn	30

List of tables

1	Comparison of multiprocessing and non-multiprocessing self-play	28
---	---	----

List of abbreviations

Glossary

1 Introduction

Chess is not only one of the most popular board games in the world, it is also a breeding ground for complex algorithms and more recently, machine learning. Chess is theoretically a deterministic game: no information is hidden from either player and every position has a calculable set of possible moves. Because the branching factor of chess is about 35-38 moves [1], calculating if a position is winning or losing requires an enormous amount of calculations.

Throughout the entire history of computer science, researchers have continuously tried to find better ways to calculate if a position is winning or losing. The most famous example is the StockFish engine [2], which uses the minimax algorithm with alpha-beta pruning to calculate the best move.

Recently, researchers at Google DeepMind have developed a new algorithm called AlphaZero [3]. This thesis explores the concept of AlphaZero, how to create a chess engine based on it, and the impact of the algorithm on both the world of chess and the rest of society.

Research has been conducted by investigating what is needed to recreate the results of AlphaZero, by programming a simple implementation using Python and TensorFlow Keras. This was done as part of a research project between November 2021 and January 2022. The code was written with lots of trial and error, as DeepMind released very little information about the detailed workings of the algorithm. It also only released a simple version of the algorithm in pseudocode, so it isn't possible to directly compare AlphaZero with other chess engines.

As I have been playing chess for most of my life, I was naturally very interested in the workings of AlphaZero and why it worked so well against the StockFish engine. Because AlphaZero was trained on supercomputers, I wanted to investigate where its flaws lie when implementing it in Python on consumer hardware.

2 Research

2.1 What is a chess engine?

According to Wikipedia [4], a chess engine is a computer program that analyzes chess or chess variant positions, and generates a move or list of moves that it regards as strongest. Given any chess position, the engine will estimate the winner of that position based on the strength of the possible future moves up to a certain depth. The strength of a chess engine is determined by the amount of moves, both in depth and breadth, that the engine can calculate.

2.2 How do traditional chess engines work?

Contemporary chess engines, like StockFish [2], use a variant of the minimax algorithm that employs alpha-beta pruning.

2.2.1 The minimax algorithm

The minimax algorithm [5] is a general algorithm usable in many applications, ranging from artificial intelligence to decision theory and game theory. The algorithm tries to minimize the maximum amount of loss. In chess, this means that the engine tries to minimize the possibility for the worst-case scenario: the opponent checkmating the player. For games where the player needs to maximize a score, the algorithm is called maximin: maximizing the minimum gain.

Minimax recursively creates a search tree [6], with chess positions as nodes and chess moves as edges between the nodes. Each node has a value that represents the strength of the position for the current player. At the start of the algorithm, the tree only consists of a root node that represents the current position. It then explores the tree in a depth-first manner by continuously choosing random legal moves, creating nodes and edges in the process.

This means that it will traverse the tree vertically until a certain depth is reached:




Figure 1: Depth-First search vs Breadth-First search [6]

When that happens, that leaf node's position is evaluated and its value is returned upwards to the parent node. The parent node looks at all of its children's values, and receives the maximum value when playing white, and the minimum value when playing black.

This repeats until the root node receives a value: the strength of the current position.

2.2.2 The evaluation function

The value estimation of leaf nodes is done by an evaluation function [7] written specifically for the game. This function can differ from engine to engine, and is usually written with help from chess grandmasters.

2.2.3 Pseudocode

The algorithm is recursive; it calls itself with different arguments, depending on which player's turn it is. In chess, white wants to maximize the score, and black wants to minimize it. [5]

```
1 function minimax(node, depth, maximizingPlayer) is
2   if depth = 0 or node is a terminal node then
3     return the heuristic value of node
4   if maximizingPlayer then
5     value := - inf
6     for each child of node do
7       value := max(value, minimax(child, depth - 1, FALSE))
8     return value
9   else (* minimizing player *)
10    value := + inf
11    for each child of node do
12      value := min(value, minimax(child, depth - 1, TRUE))
13    return value
```

Calling the function:

```
1 // origin = node to start
2 // depth = depth limit
3 // maximizingPlayer = TRUE if white, FALSE if black
4 minimax(origin, depth, TRUE)
```

2.2.4 Alpha-beta pruning

Because the necessary amount of nodes to get a good estimation of the strength of a position is so high, the algorithm needs to be optimized. Alpha-beta pruning [8] aims to reduce the amount of nodes that need to be explored by minimax. It does this by cutting off branches in the search tree that lead to worse outcomes.

Say you're playing the white pieces. You want to minimize your maximum loss, which means you want to make sure that black's score is as low as possible. Minimax always assumes that the opponent will play the best possible move. If one of white's possible moves leads to a position where black gets a big advantage, it will eliminate that branch of the search tree. As a result, the amount of nodes to explore is greatly reduced, while retaining a good estimation of the strength of the position.

2.3 Monte Carlo Tree Search

The biggest problem with minimax algorithms that use a depth limit is the dependency on the evaluation function. If the evaluation function makes incorrect or suboptimal estimations, the algorithm will suggest bad moves. Developers of contemporary chess engines like StockFish continuously try to improve this function. Since 2020, StockFish has been using a sparse and shallow neural network as its evaluation function. This neural network is still trained using supervised learning, not (deep) reinforcement learning.

Using alpha-beta pruning can also bring about some problems. Say the player can sacrifice a piece to get a huge advantage later in the game. The algorithm might cut off the branch and never explore that winning line, because it considers the sacrifice a losing position [9].

Monte Carlo Tree Search (MCTS) [10] is a search algorithm that can be used to mitigate these problems. MCTS approximates the value of a position by creating a search tree using random exploration of the most promising moves.

2.3.1 The algorithm

To create this search tree for a certain position, MCTS will run the following algorithm hundreds of times. Each of these runs is called an MCTS **simulation**. One simulation consists of four steps:

1. Selection:

- Starting from the root node, select a child node based on a formula of your choice.
- Most implementations of MCTS use some variant of Upper Confidence Bound (UCB) [11]
- Keep selecting nodes until a node has been reached that has not been visited (=expanded) before. We call this a leaf node.
- If the root node is a leaf node, we immediately proceed to the next step.

2. Expansion:

- If the selected leaf node is a terminal node (the game ends), proceed to the backpropagation step.
- If it doesn't, create a child node for every possible action that can be taken from the selected node.

3. Simulation:

- Choose a random child node that was expanded in the previous step.
- By only choosing random moves, simulate the rest of the game from that child node's position.

4. Backpropagation:

- Return the simulation's result up the tree.
- Every node tracks the number of times it has been visited, and the number of times it has lead to a win.

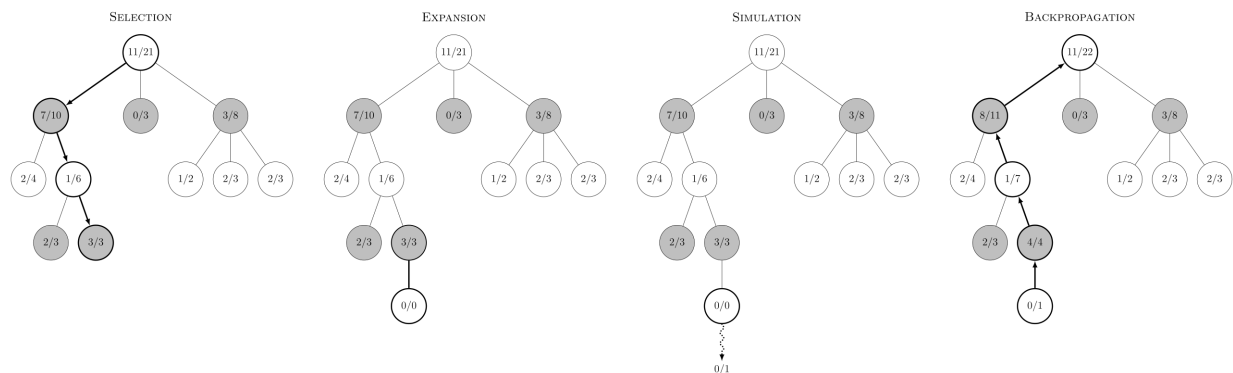


Figure 2: The 4 steps of the MCTS algorithm [10]

For chess, this algorithm is very inefficient, because of its necessity to simulate an entire game of chess in the third step of every simulation. Just to calculate the value of one position would need hundreds of these simulations to get a good estimation. This is why the selection formula needs to be chosen carefully; it's important to select nodes in a way that balances exploration and exploitation.

2.4 Go

Go is a Chinese two-player strategy board game that uses white and black stones as playing pieces [12]. It is played on a rectangular grid of (usually) 19 by 19 lines. The rules are relatively simple, but due to its

extremely complex possibilities, Go has been a very popular playground for AI research similarly to chess. Go's much larger branching factor compared to chess makes it very difficult to evaluate a position using traditional methods like minimax with alpha-beta pruning.



Figure 3: Go board [12]

2.4.1 AlphaGo

In 2014, DeepMind Technologies [13], a subsidiary of Google, started developing a new algorithm called AlphaGo to play Go [14]. Previously, the strongest Go engines were still only good enough to win against amateur Go players [15]. The algorithm used a combination of the MCTS algorithm and a deep neural network to evaluate positions.

AlphaGo was built [15], [16] by first training a neural network with supervised learning by using data from human games. The weights of that network were then copied to a new reinforcement learning network. That network was used to create a training set through self-play. By playing against itself and recording the board state, the moves the network considered, and the eventual winner of the game, a training set was created. That training set was then used to train the reinforcement learning network. A separate network (the value network) was used to estimate the value of a position.

2.4.2 AlphaGo Zero

Because AlphaGo still used some amateur games to learn from, the next step was creating a version of AlphaGo that learns completely from scratch. That's why DeepMind developed AlphaGo Zero [17]. AlphaGo Zero uses a different kind of network than AlphaGo. Instead of using two separate networks, it will combine the two outputs into a policy head and a value head. It's also using different layers: residual layers instead of convolutional layers [18].

2.5 AlphaZero

AlphaZero is a generalized version of AlphaGo Zero, created to master the games of chess, shogi ("Japanese chess"), and Go [3], [19]. For chess, AlphaZero was evaluated against StockFish version 8 by playing 1000 games with 3 hours per player, plus 15 seconds per move. It won 155 times, lost 6 times and the remaining games were drawn. AlphaZero uses a single neural network with two outputs, just like AlphaGo Zero.

2.5.1 Neural network input

The input to the network represents the state of the game. It has the following shape: $N \cdot N \cdot (M \cdot T + L)$:

- N is the board size
 - $N = 8$ in chess.
- M is the number of different pieces on the board,
 - Two players with six types of pieces each
 - Every piece is represented by its own 8x8 board of boolean values
 - For every square: 1 if the piece is on that square, 0 if it isn't
 - $M = 12$ in chess
- T is the amount of previous moves that are used as input, including the current move.
 - AlphaZero used $T = 8$ for both chess, shogi, and Go.
 - This gives the network a certain history to learn from
- L represents a set of rules specific to the game
 - $L = 7$ in chess
 - 1 plane to indicate whose turn it is
 - 1 for the total amount of moves played so far
 - 4 for castling legality (white and black can both castle kingside or queenside under certain conditions)
 - 1 to represent a repetition count (in chess, 3 repetitions results in a draw)
- $\Rightarrow 8 \cdot 8 \cdot (12 \cdot 8 + 7)$.

This means that the input to the neural network is 119 8x8 boards of boolean values.

2.5.2 Neural network layers

DeepMind tested multiple neural network architectures for AlphaGo Zero [20]. The following parts were used in these networks:

- Convolutional block
 - A convolution layer
 - Batch normalization layer
 - ReLu activation function
- Residual block
 - This consists of two convolutional blocks, and a skip-connection
 - The skip-connection will combine the input of the block to the output of the first two convolutional blocks

These networks were tested by DeepMind during development of AlphaGo Zero [18]:

1. **'dual-res'**: a single tower of 20 residual blocks with combined policy and value heads. This is the architecture used in AlphaGo Zero.
2. **'sep-res'**: two towers of 20 residual blocks each: one with the policy head and one with the value head.

3. **'dual-conv'**: a single tower of 12 convolutional blocks with combined policy and value heads.
4. **'sep-conv'**: two towers of 12 convolutional blocks each: one with the policy head and one with the value head. This is the network used in AlphaGo.

AlphaZero uses the same network architecture as AlphaGo Zero: dual-res.

2.5.3 Neural network output

The neural network has two outputs:

- A policy head, which represents a probability distribution over the possible actions.
- A value head, which represents the value of the current position.

While the value head simply outputs a single float value between -1 and 1, the policy head is quite a bit more complicated. It outputs a vector of probabilities, one for each possible action in the chosen game. For chess, 73 different types of actions are possible:

- 56 possible types of “queen-like” moves: 8 directions to move the piece a distance between 1 and 7 squares.
- 8 possible knight moves
- 9 special “underpromotion” moves:
 - If a pawn is promoted to a queen, it is counted as a queen-like move (see above)
 - If a pawn is promoted to a rook, bishop, or knight, it is seen as an underpromotion (3 pieces)
 - 3 ways to promote: pushing the pawn up to the final rank, or diagonally taking a piece and landing on the final rank
 - $\Rightarrow 3 \cdot 3 = 9$

These 73 actions are each represented by a plane of 8×8 float values. Say the first plane is a queen-like move to move a piece one square northwest, the second plane could be the same type of move, but a distance of two squares, and so on. The squares on these planes represent the square from which to pick up a piece.

The result is a $73 \cdot 8 \cdot 8$ vector of probabilities, so 4672 float values.

2.6 Training the network

To train this type of network, it's necessary to create a large dataset. This is done by letting the engine play against itself for a high amount of matches. For every action taken by the agent, data is collected and stored in the training set. For complex games like chess, shogi and Go, this training set needs to be huge because of the extremely large amount of possible situations.

2.6.1 Tensor Processing Units (TPU)

Because of the requirement to play a high amount of matches against itself, it was necessary to calculate MCTS simulations in parallel on as fast as possible hardware. To help with these calculations, DeepMind used Google's newly created Tensor Processing Units (TPU) [21]. A TPU is an application-specific integrated circuit (ASIC [22]) that is specifically built for machine learning with neural networks. Since 2018, these TPUs have been made publicly available to rent through Google's Cloud Platform. Smaller TPUs can be purchased from Google.

2.7 Leela Chess Zero

Leela Chess Zero (lc0) is a free, open-source project that attempts to replicate the results of AlphaZero [23]. Lc0 was adapted from Leela Zero [24], a Go computer that attempted to replicate the results AlphaGo Zero [17].

It is written in C++ [25], and it has managed to play at a level that is comparable to the current best version of StockFish. Because lc0 is a community driven project, volunteers can help create training games through self-play using their own computers. This made it possible to feed millions of chess games into the network.

3 Technical research

3.1 Introduction

Creating the chess engine required programming the following parts:

- The MCTS algorithm
- A tree data structure with nodes and edges
- The neural network
- The training pipeline
- The evaluation pipeline
- A way to store every move to a dataset
- A class to make the engine play against itself
- A GUI to play against the engine
- Docker containers for easily scaling and distributing the program

All code was written in Python. The neural network was made using the TensorFlow Keras library. Chess rules and helper functions were implemented using the open-source library python-chess [26].

3.2 Class structure

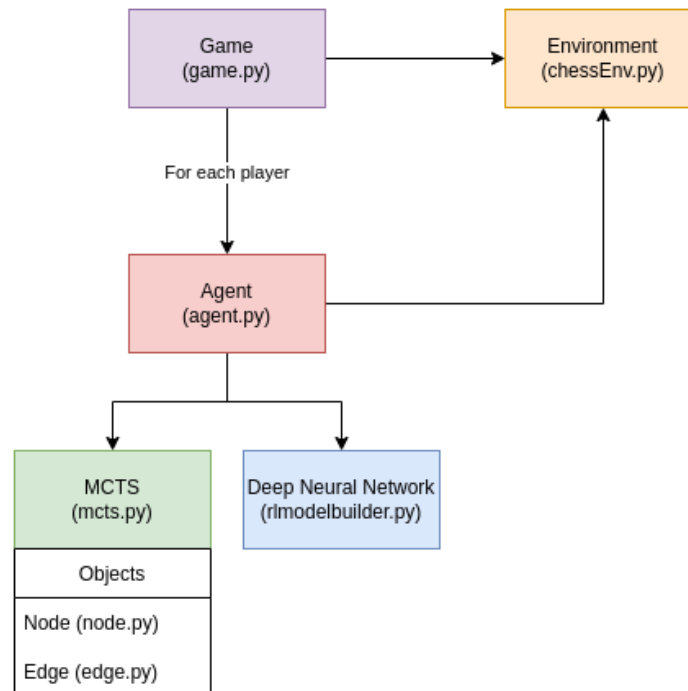


Figure 4: Basic class structure for the code responsible for playing a game

3.2.1 Making one move

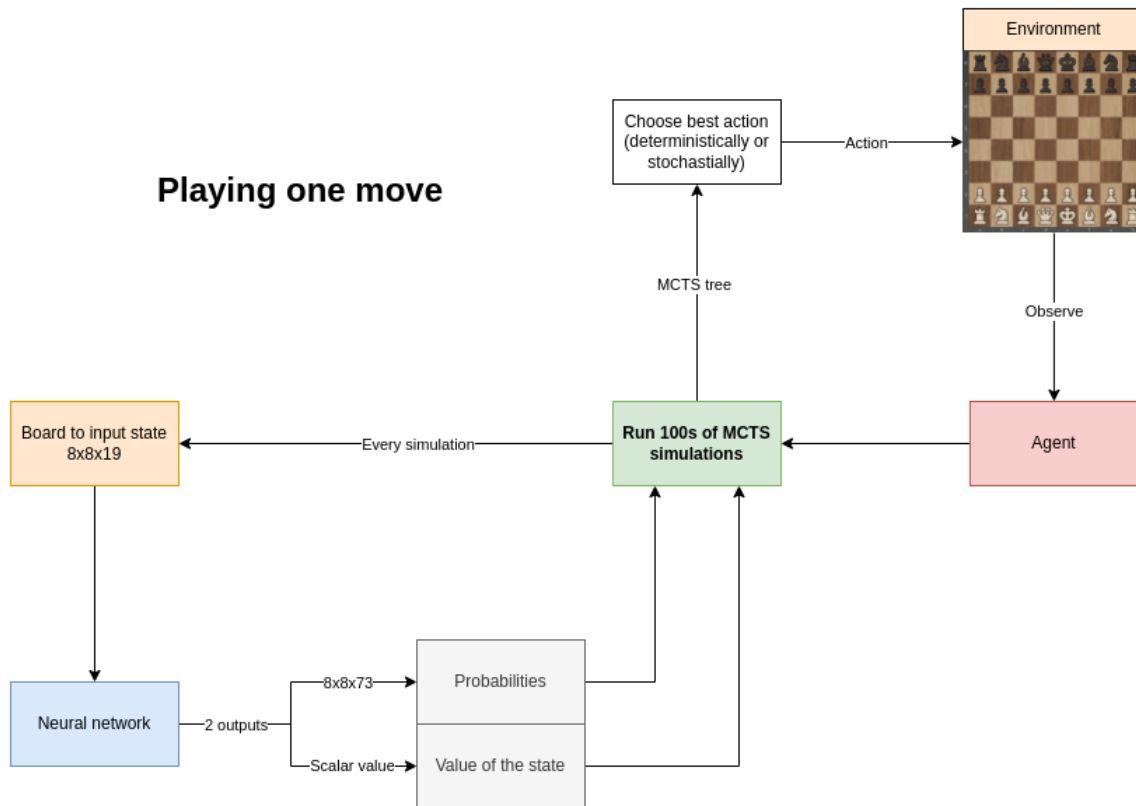


Figure 5: Pipeline to make one move

To make a move, an Agent (white or black) observes the environment: the chessboard. The agent calls upon the MCTS class to create a tree with as root the current state of the chessboard. The MCTS class will run the MCTS algorithm hundreds of times. This amount is configurable in the config file. Higher amounts result in a more accurate estimation of the position's value, but also in longer computation times.

Every MCTS simulation, the neural network will be called to evaluate a position. The two outputs, the policy and the value, will be used to update the tree.




Figure 6: Example tree after 400 simulations. N = amount of times the selection step selects that edge

After the simulations are done, the agent will pick the best move from the tree. It can do this in two ways:

- Deterministically: choosing the most visited move
- Stochastically: creating a uniform distribution of the visit counts and picking a move from that distribution

Stochastic selection is better when creating a training set, as it will result in a more diverse dataset. Deterministic selection is better when evaluating with a previous network, or playing against the network competitively. In that case, picking the most visited move is the best choice.

3.3 The neural network

Initially, a prototype of the neural network was created with randomly initialized weights. The Python class to create the model was immediately made with customizability in mind: the input and output shapes can be given as arguments, and the sizes of the convolution filters can be changed using a configuration file.

The neural network architecture is the same as AlphaZero's (see the Research section).

3.4 A tree structure with nodes and edges

As mentioned before, the MCTS algorithm creates a tree structure to represent the possible future states after the current position.

Node and Edge classes were written. The Node class represents a position in the game, and holds the following data:

- The position: a string representation of the board using the Forsyth-Edwards Notation (FEN) [27]
- The current player to move (boolean)
- A list of edges connected to this node
- The visit count of this node, initialized to 0
- The value for this node, initialized to 0

The Edge class represents a move. It holds the following data:

- The input node (the position from which the move was made)
- The output node (the resulting position after taking the move)
- The move itself: an object of the Move class from the python-chess library, which holds:

- The source square and the target square of the move
- If the move was a promotion: the piece it was promoted to
- The prior probability of this move
- The visit count of this edge, initialized to 0
- The value for this action, initialized to 0

The tree can also be plotted using the Graphviz library [28]. A recursive function was written to create the tree and output it to an SVG file.

3.5 The MCTS algorithm

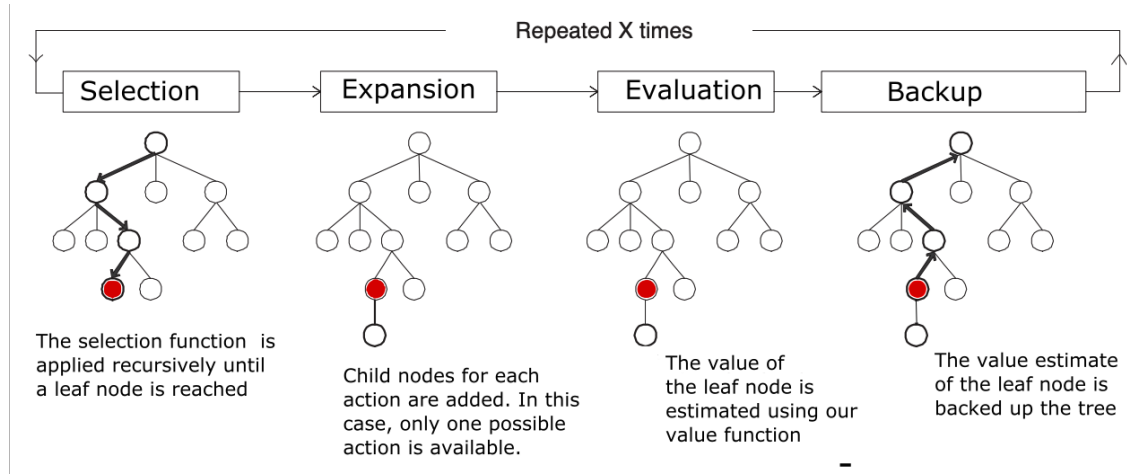


Figure 7: The four steps in AlphaZero's MCTS algorithm [29]

A class was written to hold an MCTS tree. It holds the agent who created the tree and the current position as the root node.

3.5.1 The selection step

For the selection step, the following UCB formula [19] was used to determine which edge to select:

$$UCB = \left(\log\left(\frac{1 + N_{\text{parent}} + C_{\text{base}}}{C_{\text{base}}}\right) + C_{\text{init}} \right) \cdot P \cdot \frac{\sqrt{N_{\text{parent}}}}{(1 + N)} \quad (1)$$

- C_{base} and C_{init} are constants that can be changed in the config file. The same values as AlphaZero were used.
- N_{parent} is the visit count of the input node
- N is the visit count of the edge
- P is the prior probability of the edge

The selection step combines this UCB formula with the edges action-value and visit count:

$$Q = \frac{W}{N + 1} \quad (2)$$

$$V = \begin{cases} UCB + Q & \text{if white} \\ UCB - Q & \text{if black} \end{cases} \quad (3)$$

The edge with the highest value (V) is selected. After selection, the visit count for the edge's output node is incremented by one. We call this output node the 'leaf node'.

3.5.2 The expansion step

The leaf node is expanded by creating a new edge for each possible (legal) move. If there are no legal moves, the outcome (draw, win, loss) is checked and the leaf node is passed to the next step in the algorithm.

If there are legal moves, the leaf node is given as an input to the neural network.

While AlphaZero uses an input shape of 119 8x8 boolean boards, I opted for a much lighter input shape. The input used in this project only has 19 boards:

- 1 board to show which turn it is: 1 is white and 0 is black
- 4 boards to show if each player still has castling rights
- 1 board to show if a draw is possible after 50 moves without a capture or pawn move
- 12 boards to show the positions of the pieces on the board
- 1 board to show the en passant square: if a pawn can be taken en passant, this square is set to 1

For example:



Figure 8: Board example



Figure 9: The input state converted from the above example board

Notice how the last square shows the square where en passant is possible. Grey padding was added to make a better visual presentation of the output planes.

The neural network will return the policy and the value of the position. As described in the research part of this thesis, the value is a float between -1 and 1, and the policy is a 73x8x8 tensor of floats. This policy is then mapped to a dictionary, where keys are moves and the values are the probabilities of each move.

The neural network gets no information about the rules of chess, so it is necessary to filter out the illegal moves.



Figure 10: A subset of the 73 8x8 planes from the policy output. The brighter the pixel, the better the move.



Figure 11: The same subset, but with the illegal moves filtered out

The above two images were made using a trained model. The policy output of an untrained model with random weights looks like this:




Figure 12: Some output planes from an untrained model.

(Black padding was used instead of gray to make it clearer.)

This clearly shows the trained model has at least some level of understanding of which moves are legal, without giving it any knowledge about the rules of chess. The model isn't even told we're playing chess: it learns the correct outputs completely on its own.

3.5.3 The evaluation step

The value received from the neural network is now assigned to the leaf node.

3.5.4 The backpropagation step

The value from the leaf node is now also added to every selected node in the path from the root to the leaf node. Concretely, for every selected edge in the path, the following values are changed:

- The edge's input node's visit count is incremented by 1
- The edge's visit count is incremented by 1
- The edge's value is incremented by the value of the leaf node

3.6 Creating the dataset




Figure 13: The training set consists of: the input state, the move probabilities, and the eventual winner

To improve the performance of the neural network, the algorithm needs to play against itself for a high amount of games. Every move is saved to memory in a simple Python list. Once the game is over, the winner is added to every move in memory. The whole game is then saved to a binary file in NumPy's .npy format. When a new game starts, the memory is cleared.

3.6.1 Creating a dataset from puzzles

Because creating the dataset is extremely time-consuming, I came up with the idea to also create data from chess puzzles. The idea is to let the agent play from a certain given position (the puzzle), instead of from the start of the game, and see if the agent can find the correct solution. This would in theory give the model a better understanding of common chess tactics, without having to play whole games from start to finish.

A chess puzzle is a position with one simple goal: find the best move or sequence of moves. Lichess.org, the open source chess website, has a huge database of over 2 million of these puzzles publicly available for free [30]. These are the most common puzzle categories:

- Mate-in-X: find the best moves to checkmate the opponent in a maximum of X moves
- Capturing one of the opponent's undefended pieces
- Getting a positional advantage

When creating the dataset, mate-in-1 and mate-in-2 puzzles were also added to the dataset. This is because these kinds of puzzles end quickly (and are thus less time-consuming), and they have a clear winner: the player who can checkmate the opponent. The Lichess database has over 400,000 puzzles of these two types.

3.7 Training the neural network

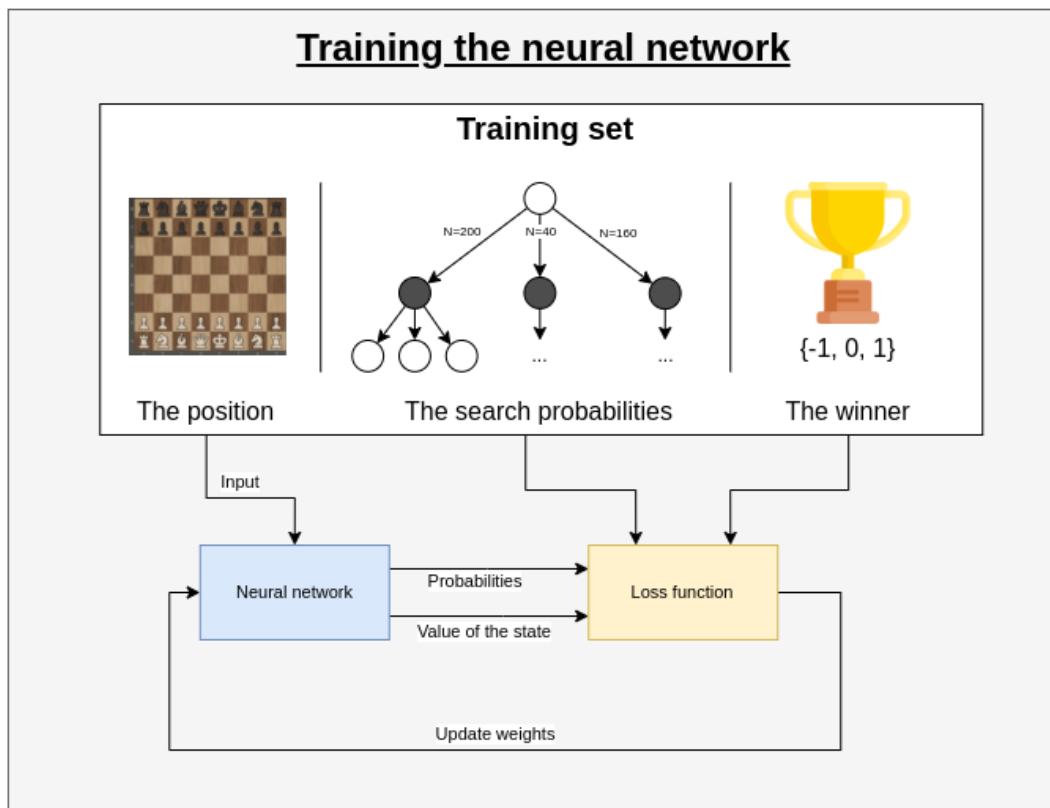


Figure 14: The training pipeline

To train the neural network, every saved game's binary file is loaded into memory. The memory is shuffled to avoid the neural network accidentally learning time dependent patterns. Random batches of the training set are then processed through the neural network.

The position is used as the input to the neural network. The network's outputs are then compared to the move probabilities and the winner from the dataset:

- The move probabilities are converted to a $73 \times 8 \times 8$ tensor, which is compared with the output of the network
- The winner is compared with the value output of the network.

The training pipeline employs two separate loss functions. The policy head uses categorical cross-entropy and the value head uses mean squared error.

3.7.1 The first training session

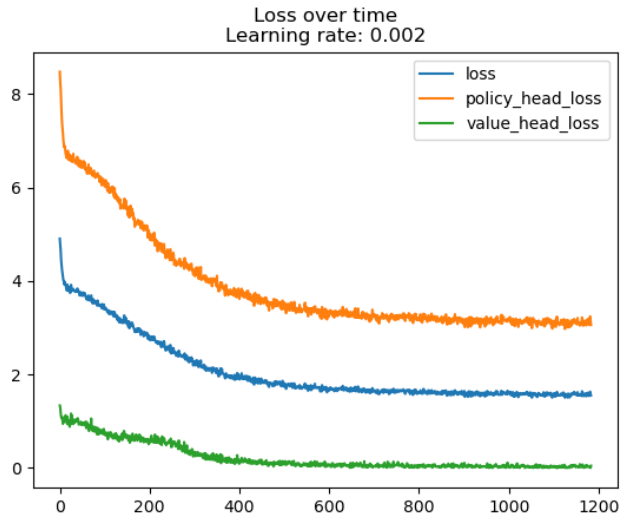


Figure 15: First training session

Naturally, the first training session started with a random model. The dataset was created by running self-play for many hours, resulting in a dataset of around 76,000 positions. Training was performed with a learning rate of 0.002, and a batch size of 64. The model uses the Adam optimizer.

After training, the model was saved to disk.

3.7.2 The second training session

The previous model was used as the starting point for the second training session. A new dataset was made using only games created by that model. This time, due to time constraints, the dataset only had 50,000 positions.




Figure 16: Second training session

It's clear the model manages to learn something at the start of the training session, but seems to plateau after 150 batches. Because the first training session started from a random model, it managed to learn a lot more before reaching a plateau after around 600 batches.

3.8 Multiprocessing

Because self-play is very time-consuming, there needed to be a way to play multiple games in parallel.

3.8.1 Without multiprocessing




Figure 17: Self-play without multiprocessing

Previously, every game was played in its own process, but every agent needed to send predictions to a neural network. This resulted in every process creating its own copy of the neural network. This is extremely heavy for the GPU, and it's impossible to scale. Due to VRAM limits, only two games could be played in parallel on my system (RTX 3070).

3.8.2 With multiprocessing

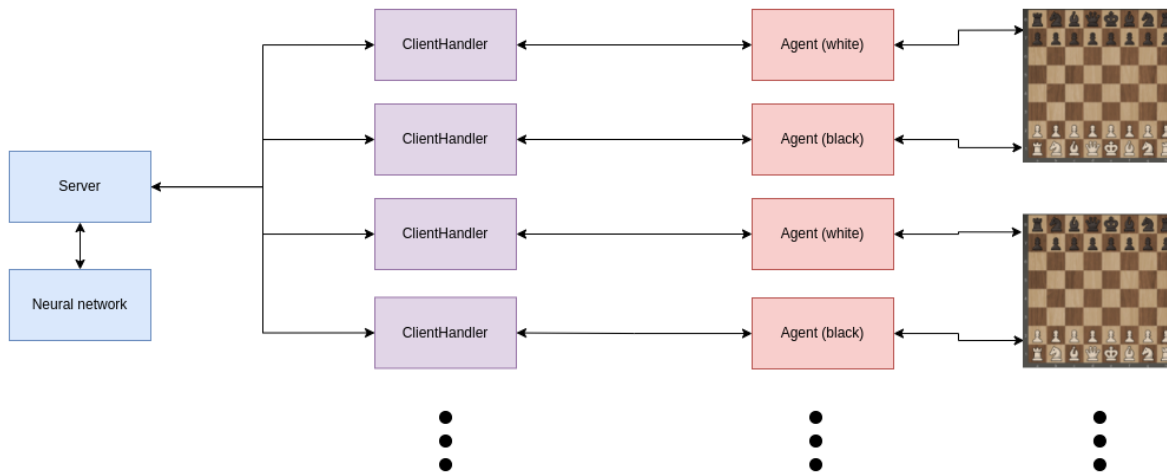


Figure 18: Self-play with multiprocessing

To solve the issue of parallel self-play, I created a client-server architecture with Python sockets. The server side has access to the neural network. For every client (the self-play process), the server creates a ClientHandler in a new thread.

Here's how one client works:

1. The MCTS algorithm runs 100s of simulations every move. Every simulation sends a chess position to the neural network.
2. The agent sends the chess position through a socket.
3. The ClientHandler receives the chess position and sends it to the server.
4. The server calls the network's predict function and returns the outputs to the client handler.
5. The ClientHandler sends the outputs back to the client.

This socket communication does have a small overhead in both time and CPU usage, but it's much faster than the previous method because it is much more scalable. Here's a comparison of the two methods, with the server running on the first system:

	No multiprocessing	With multiprocessing (8 games)
RTX 3070 + Ryzen 7 5800H	50 simulations/sec	30 simulations/sec per game
GTX 1050 + i7 7700HQ	30 simulations/sec	15 simulations/sec per game

Table 1: Comparison of multiprocessing and non-multiprocessing self-play

With 8 games in parallel on the first system, I managed to get an average speed of $30 \cdot 8 = 240$ simulations per second. The second system is much less powerful and needed to connect over Wi-Fi, but it still managed an average of 120 simulations per second.

3.9 A GUI to play against the engine

The GUI was based on a GitHub project that someone created for simply visualizing a chessboard with PyGame [31], [32]. I greatly improved and expanded that project to include a way for the player to interact

with the board, play against an engine, play against yourself, and more [33]. I created a pull request to include my changes in the original author's GitHub project, but it was declined due to being too extensive and out-of-scope for the project [34].

Using this code, a GUI class was created for playing against the engine. The player can choose a color and play against the engine. Clicking on a piece will highlight its square, clicking on another square will move the piece to that square. Right-clicking will cancel the move.



Figure 19: Chessboard example when playing against the engine

Promoting a pawn can be done by simply moving a pawn to the last rank. A menu will pop up to choose the piece to promote to.



Figure 20: Promoting a pawn

3.10 Docker images

Because the amount of data needed to train the neural network is very large, two docker images were created:

- A server image that runs the neural network and listens for clients.
- A client image that runs self-play and sends chess positions to the server.

With a docker-compose file, the server and client images can be easily started. The number of clients to run in parallel can be easily configured in that file.

It's recommended to run the server on a fast GPU-equipped system, and the clients on a system with many high-performance CPU cores. The clients do not need a GPU to run self-play.

3.11 The final project

This section walks through all executable programs included in the project [35].

3.11.1 Creating your own untrained AI model

```

1  $ python rlmodelbuilder.py --help
2  usage: rlmodelbuilder.py [-h] [--model-folder MODEL_FOLDER] [--model-name MODEL_NAME]
3
4  Create the neural network for chess
5
6  options:
7  -h, --help            show this help message and exit
8  --model-folder MODEL_FOLDER
9                        Folder to save the model
10 --model-name MODEL_NAME
11                        Name of the model (without extension)

```


This will create a new model with the name <NAME> in the folder <FOLDER>. The model parameters (amount of hidden layers, input and output shapes if you want to use the network for a different game, the amount of convolution filters, etc.) can be changed by editing the config.py file.

3.11.2 Creating a training set through self-play

Creating the docker containers:

```
1 # in the repo's code/ folder:
2 docker-compose up --build
```

This will create one server and the amount of clients that is configured in the docker-compose.yml file. That file can also be used as a reference to deploy a Kubernetes cluster for parallel self-play with high scalability and reliability.

Using the created GUI, it is possible to visualize the self-play of all these boards in real-time. This can be done by settings the “SELFPLAY_SHOW_BOARD” environment variable in docker-compose.yml to “true”. For every replica of the client, a new PyGame window will open where the board will change in real-time.

You can also manually run self-play or create data using puzzles:

```
1 $ python selfplay.py --help
2 usage: selfplay.py [-h] [--type {selfplay,puzzles}] [--puzzle-file PUZZLE_FILE]
3       [--puzzle-type PUZZLE_TYPE] [--local-predictions]
4
5 Run self-play or puzzle solver
6
7 options:
8   -h, --help            show this help message and exit
9   --type {selfplay,puzzles}
10                        selfplay or puzzles
11   --puzzle-file PUZZLE_FILE
12                        File to load puzzles from (csv)
13   --puzzle-type PUZZLE_TYPE
14                        Type of puzzles to solve. Make sure to set a
15                        puzzle move limit in config.py if necessary
16   --local-predictions    Use local predictions instead of the server
```

3.11.3 Evaluating two models

To determine whether your new model is better than the previous best, you can use the evaluate.py script. It will simulate matches between the two models and record the wins, draws and losses.

In chess, white inherently has a slightly higher chance of winning because they can play the first move [36]. Therefore, to evaluate two models, each model will both play white and black an equal amount of times.

```
1 $ python evaluate.py --help
2 usage: evaluate.py [-h] model_1 model_2 nr_games
3
4 Evaluate two models
5
6 positional arguments:
7   model_1    Path to model 1
8   model_2    Path to model 2
9   nr_games    Number of games to play (x2: every model plays both white and black)
```

```
10
11 options:
12   -h, --help  show this help message and exit
```

3.11.4 Playing against the AI

```
1 $ python main.py --help
2 usage: main.py [-h] [--player {white,black}] [--local-predictions] [--model MODEL]
3
4 options:
5   -h, --help            show this help message and exit
6   --player {white,black}
7                           Whether to play as white or black. No argument means random.
8   --local-predictions    Use local predictions instead of the server
9   --model MODEL          For local predictions: specify the path to the model to use.
```

This will start the GUI application to allow you to play against the engine.

3.12 Porting to C++

Optimization is extremely important for chess engines, especially when the engine needs to play against itself to create a dataset. That is why Python was not an ideal choice to implement this chess engine.

Currently, I'm writing a C++ version of the engine in my spare time [37]. Instead of TensorFlow Keras, I've opted to use PyTorch instead. So far, I've noticed that the C++ version with PyTorch manages to run the MCTS algorithm four times faster than the Python version: 200 simulations per second instead of 50.

4 Reflection

4.1 Introduction

This section is a reflection on the project. It will answer the following questions:

- What are the strengths and weaknesses of this research project?
- Is the result of the project usable for corporations?
- What are the possible obstacles for companies that wish to implement this?
- What is the added value for companies?
- Which alternatives are there?
- Is there a socio-economic impact present?
- Is there opportunity for further research?

To help answer these questions, I have consulted with Dr. Thomas Moerland, a postdoctoral researcher at the Reinforcement Learning Group at Leiden University, in The Netherlands [38]. With his experience in researching AI and specifically Reinforcement Learning, he was able to give me some insights into how AlphaZero could impact society in industries besides gaming.

I have also asked questions to the following online communities:

- The AI Stack Exchange community
- The Leela Chess Zero Discord community
- The TalkChess forum, a forum affiliated with The Computer Chess Club (CCC)

4.2 Strengths and weaknesses

This project successfully implements the MCTS algorithm as modified by DeepMind for AlphaZero and their preceding work, AlphaGo and AlphaGo Zero. The resulting chess engine can be played against using the GUI, and it can be used to create a dataset for further improving the neural network. The project serves as an example of how AlphaZero's MCTS algorithm can be implemented in a chess engine using Python and consumer-grade hardware.

The project offers a way to easily create your own neural network with the `rlmodelbuilder.py` script. The code is written with modularity and adaptability in mind: by changing parameters in the `config.py` file, you can easily create a different model for chess, or for similar applications with different input and output shapes. The MCTS algorithm can be used as a reference if you want to implement it other applications.

The ultimate goal of this research project was to create a chess engine that is at least capable of winning against amateur players of around 1000 ELO. This was not possible with my consumer-grade hardware: the chess engine still seems to play very randomly and has not found a good winning strategy. The biggest hurdle was the lack of computational resources to create a large dataset.

However, the project does offer a solution to create more data through self-play: with the docker containers, it is easy for other people to help create a dataset by deploying their own cluster of servers and clients. This way, a global dataset could be made by combining self-play data from volunteers around the world. It can then be used to train a new model.

This is exactly how Leela Chess Zero operates: the developers created a simple program for volunteers to run, which makes the latest model play against itself and uploads the data back to the developers. Once there is enough data, the developers can retrain the model. By continuously repeating this process, Leela Chess Zero slowly became a powerful chess engine on par with the best chess engines in the world [39].

4.3 Is the result of the project usable in the corporate world?

While AlphaZero and its variants have mostly been used in gaming applications, the project is not limited to this. For instance, AlphaGo has been used by Google to properly ventilate their data centers more efficiently. The algorithm helped lower data center cooling costs by 40% [40]. It does this by periodically pulling data from many sensors and feeding that data to a deep neural network to predict how different combinations of potential actions will affect future energy consumption [41].

Theoretically, this project can be used to solve virtually any problem that can be defined as an agent acting upon an environment with discrete state and action spaces [42].

Reinforcement learning in general has been used to solve a wide range of problems. Many websites that have recommendation systems have already used reinforcement learning to improve which products get recommended to users [43]. For example, Netflix uses contextual bandits to personalize artwork to the user [44], [45]. This avoids the need for waiting to collect data from users, training a model, waiting for an A/B test to conclude, to then finally recommend the best artwork. By then, too many users have ignored an item they would have otherwise clicked on.

In the field of robotics, (deep) reinforcement learning can be used to teach a robot through self-play to perform simple tasks that previously required manual programming.

4.4 What are the possible obstacles for companies that wish to implement this?

Depending on the application, the amount of data necessary to train a sufficiently powerful model can be extremely large. It's necessary for the company to have a good infrastructure to create this data. Companies can also invest in cloud solutions like the Google Cloud Platform to take advantage of their TPUs [46].

Because these types of algorithms need to be trained through self-play, applications should be written in a programming language that is capable of training and inferring neural networks, but should also be very performant and efficient. Fast, low-level programming languages like C++ often require much more development time to create the same application compared to high-level programming languages like Python. This should also be factored into the development process.

AI frameworks like TensorFlow and PyTorch are available in both Python and C++, but documentation and support is lacking in the C++ versions. This means companies who wish to implement this type of application will need to hire C++ developers who are also experienced with neural networks.

4.5 What is the added value for companies?

Options are always useful, and this project offers a way to solve problems using AI without having to manually collect data. The reinforcement learning agent will collect data for you by playing against itself. This means the algorithm will get better in a "fire-and-forget" manner: just launch a cluster of self-play agents and wait until it gathers enough data. You could create an automatic pipeline to retrain the model after a set amount of time or after a certain amount of games.

4.6 Which alternatives are there?

For chess specifically, there are many chess engines that can be used for both analysis and playing against. StockFish is the most popular, and as mentioned before, Leela Chess Zero is also open source and works the same way as AlphaZero.

For industries other than the gaming industry, many problems are solvable by using machine learning algorithms instead of reinforcement learning solutions. Supervised learning methods have a very wide range of applications, provided you have the data to train a model.

However, if you don't have a means to gather data, there are many reinforcement learning algorithms other than AlphaZero:

- Q-learning

- SARSA
- DQN
- Actor-Critic methods like A2C or A3C
- ...

Currently, sending inputs from the client to the server for predicting outputs with the neural network happens through Python sockets. A good alternative for that is gRPC. It would allow the clients to directly call a method on the server without having to resort to Python sockets, which might be slower. [47]

4.7 Is there a socio-economic impact present?

Sadly, algorithms like AlphaZero fail at solving even relatively simple problems like chess when the model does not have enough data to train on. [38]. Given enough time and data, though, these types of algorithms can be used to achieve superhuman performance in applications that were previously impossible to find optimal solutions for.

DeepMind's ultimate goal is to “[...] solve intelligence, developing more general and capable problem-solving systems, known as artificial general intelligence (AGI).” [48]. This is a slow and extremely difficult process, but DeepMind is confident AlphaZero and its variants are a big step towards that goal [49].

Others are more pessimistic about the progress DeepMind has made **emailmoerlandEmailExchangeTuur22**, [50], explaining that while AlphaGo (Zero) and AlphaZero have certainly been a huge breakthrough in their respective games, this does not mean their results a particularly big step towards AGI.

I think there is a huge gap between the simple games we now solve and the 'artificial general intelligence' that DeepMind is trying to achieve. I believe that the most progress won't come from algorithms, but from calculations. For general AI, we need incredibly strong simulators where a huge amount of data can be passed through for a very long time. Just like the millions of years of evolution our brains have been through, our brains have been gathering data non-stop throughout our lives. [51]

4.8 Is there opportunity for further research?

While reinforcement learning has existed for a long time, deep reinforcement learning techniques have only recently been used. This is because of the lack of computational resources to train a well-performing model. With high performance computing devices and cloud solutions getting cheaper and more powerful, a lot of research opportunities open up. Deep Reinforcement learning for autonomous driving has recently become an active area of research in both academia and industry [52].

For this project specifically, there is still a lot of research opportunity to try to solve the problem of creating a large enough dataset to train a sufficiently powerful model.

5 Advice

5.1 Introduction

This section gives advice to people who wish to implement something similar to the project. It includes recommendations and tips for when people want to program an application that uses the algorithms and concepts researched in this thesis.

5.2 When should you use this technology?

This algorithm has a very specific purpose. It should only be used in situations where both the action space and state space are discretely defined. Both of these should also be relatively small, as bigger action spaces and state spaces will require much more data to successfully train a model. Therefore, it is important to keep the problem as simple as possible.

5.3 Recommendations

When you want to implement AlphaZero or similar algorithms, it is recommended to prioritize writing code in the most efficient way possible. Every step of the algorithm should be programmed with multiprocessing and multithreading in mind. Furthermore, the choice of programming language should be considered carefully. While most AI frameworks are primarily written for use in Python, and most documentation is written for Python, the language just isn't fast enough for something that needs this much data through self-play. C++ would be a better choice.

As was apparent from the results of the research project, generating training data takes a very long time. A good solution to this problem is to first pretrain the model on a small dataset. This will result in a model that is at least slightly familiar with the environment. Without a pretrained model, the algorithm will start from nothing and make completely random actions. With a pretrained model, the algorithm will have some idea on what decisions to make based on the data it was given in pretraining.

5.4 Tips when programming a similar application

This section will give some useful tips for programming an application that uses AlphaZero's MCTS algorithm.

5.4.1 Creating the tree data structure

The MCTS algorithm requires a tree to be constructed. As mentioned before, the nodes of the tree are the positions of the game, and the edges between those nodes are the actions that can be taken from those positions.

It is important to avoid memory management issues when creating the tree. The MCTS algorithm can create quite large trees, especially when using hundreds of simulations. Every node object should therefore only include the necessary data. For example, a very early version of my MCTS algorithm included all previous moves of the game in every node, creating a huge memory leak. This was fixed by only storing the current board state as a FEN string in each node.

5.4.2 Programming the MCTS algorithm

When programming the MCTS algorithm, it is crucial to make every step as efficient as possible. In the expansion step of the algorithm, the state of the leaf node should be converted to an object that can be used as an input in the neural network. The function that converts the leaf's state can be very slow if it's not programmed efficiently. As that function needs to be called for every simulation, every line of code should be optimized for speed.

Calling the neural network is definitely the slowest step of the algorithm. A good way to take advantage of this is to send batches of data to the network instead of one input at a time. This can be done by creating

multiple threads to “crawl” through the tree at the same time during the selection step of the algorithm. Once every thread has finished by finding a leaf node, every leaf node can be sent to the network at once. For every output received by the neural network, a thread can again be started for the third and fourth steps. This will result in a significant speedup. The difficulty of programming this is mainly making sure the “crawlers” during the selection step don’t all select the same nodes. That would result in a less explored MCTS tree.

5.4.3 Choosing the right input for the neural network

As mentioned in the technical research part of this thesis, I opted for a smaller input shape than the one defined in AlphaZero’s paper. This is because I wanted to make the neural network as fast as possible.

You should carefully consider what data the neural network would need to make a decision, based on the problem you’re trying to solve. For example: when Google used AlphaGo to make their data center’s temperature control more efficient, they opted to use the sensors in the data center as inputs to the neural network. You should make sure to omit any unnecessary data from the input.

6 Conclusion

This thesis was created to critically evaluate my research project that tried to answer the question “How to create a chess engine using deep reinforcement learning”. The aim of the project was to create a chess engine based on DeepMind’s AlphaZero in Python with TensorFlow Keras. This was attempted by programming the MCTS algorithm and a neural network that can train on data collected from self-play. The ultimate goal of the project was not to create a chess engine that was capable of beating AlphaZero, but rather to create a simple version of AlphaZero that could beat amateur chess players of around 1000 ELO.

Sadly, the resulting chess engine has not managed to reach this goal. This is because of the lack of computational resources necessary to create a big enough dataset through self-play. However, this thesis does offer some solutions to this problem:

- Pretrain the model on a small dataset
- Use docker containers to set up a cluster of self-play agents

This thesis also serves as a warning to people who wish to implement similar algorithms without having the necessary access to powerful hardware. It recommends investigating certain alternatives to AlphaZero, such as other reinforcement techniques like Q-learning, SARSA, and DQN, but also supervised learning techniques if training data is available.

7 Bibliography

- [1] *Branching Factor - Chessprogramming wiki*. [Online]. Available: https://www.chessprogramming.org/Branching_Factor (visited on 03/28/2022).
- [2] "Stockfish (chess)," *Wikipedia*, Mar. 2022. [Online]. Available: [https://en.wikipedia.org/w/index.php?title=Stockfish_\(chess\)&oldid=1079146589](https://en.wikipedia.org/w/index.php?title=Stockfish_(chess)&oldid=1079146589) (visited on 03/28/2022).
- [3] "AlphaZero," *Wikipedia*, Jan. 2022. [Online]. Available: <https://en.wikipedia.org/w/index.php?title=AlphaZero&oldid=1065791194> (visited on 02/01/2022).
- [4] "Chess engine," *Wikipedia*, Apr. 2022. [Online]. Available: https://en.wikipedia.org/w/index.php?title=Chess_engine&oldid=1080874516 (visited on 04/05/2022).
- [5] "Minimax," *Wikipedia*, Mar. 2022. [Online]. Available: <https://en.wikipedia.org/w/index.php?title=Minimax&oldid=1076761456> (visited on 04/05/2022).
- [6] M. Eppes, *How a Computerized Chess Opponent "Thinks" — The Minimax Algorithm*, Oct. 2019. [Online]. Available: <https://towardsdatascience.com/how-a-chess-playing-computer-thinks-about-its-next-move-8f028bd0e7b1> (visited on 04/05/2022).
- [7] "Evaluation function," *Wikipedia*, Mar. 2022. [Online]. Available: https://en.wikipedia.org/w/index.php?title=Evaluation_function&oldid=1079533564 (visited on 04/06/2022).
- [8] "Alpha-beta pruning," *Wikipedia*, Jan. 2022. [Online]. Available: https://en.wikipedia.org/w/index.php?title=Alpha%E2%80%93beta_pruning&oldid=1068746141 (visited on 02/01/2022).
- [9] *Minimax and Monte Carlo Tree Search - Philipp Muens*. [Online]. Available: <https://philippmuens.com/minimax-and-mcts> (visited on 04/06/2022).
- [10] "Monte Carlo tree search," *Wikipedia*, Apr. 2022. [Online]. Available: https://en.wikipedia.org/w/index.php?title=Monte_Carlo_tree_search&oldid=1081107255 (visited on 04/06/2022).
- [11] *ML / Monte Carlo Tree Search (MCTS)*, Jan. 2019. [Online]. Available: <https://www.geeksforgeeks.org/ml-monte-carlo-tree-search-mcts/> (visited on 04/07/2022).
- [12] "Go (game)," *Wikipedia*, Mar. 2022. [Online]. Available: [https://en.wikipedia.org/w/index.php?title=Go_\(game\)&oldid=1079941654](https://en.wikipedia.org/w/index.php?title=Go_(game)&oldid=1079941654) (visited on 04/06/2022).
- [13] "DeepMind," *Wikipedia*, Feb. 2022. [Online]. Available: <https://en.wikipedia.org/w/index.php?title=DeepMind&oldid=1072182749> (visited on 04/07/2022).
- [14] "AlphaGo," *Wikipedia*, Mar. 2022. [Online]. Available: <https://en.wikipedia.org/w/index.php?title=AlphaGo&oldid=1077595428> (visited on 04/07/2022).
- [15] *AlphaGo*. [Online]. Available: <https://www.deepmind.com/research/highlighted-research/alphago> (visited on 04/07/2022).
- [16] *Mastering the game of Go with Deep Neural Networks & Tree Search*. [Online]. Available: <https://www.deepmind.com/publications/mastering-the-game-of-go-with-deep-neural-networks-tree-search> (visited on 04/07/2022).
- [17] "AlphaGo Zero," *Wikipedia*, Feb. 2022. [Online]. Available: https://en.wikipedia.org/w/index.php?title=AlphaGo_Zero&oldid=1073216893 (visited on 04/08/2022).
- [18] *Mastering the game of Go without Human Knowledge*. [Online]. Available: <https://www.deepmind.com/publications/mastering-the-game-of-go-without-human-knowledge> (visited on 04/07/2022).
- [19] D. Silver, T. Hubert, J. Schrittwieser, *et al.*, "Mastering Chess and Shogi by Self-Play with a General Reinforcement Learning Algorithm," *arXiv:1712.01815 [cs]*, Dec. 2017. arXiv: 1712.01815 [cs]. [Online]. Available: <http://arxiv.org/abs/1712.01815> (visited on 02/01/2022).
- [20] *Neural Networks - Chessprogramming wiki*. [Online]. Available: https://www.chessprogramming.org/Neural_Networks#ANNs (visited on 04/07/2022).
- [21] "Tensor Processing Unit," *Wikipedia*, Mar. 2022. [Online]. Available: https://en.wikipedia.org/w/index.php?title=Tensor_Processing_Unit&oldid=1077688658 (visited on 04/07/2022).
- [22] "Application-specific integrated circuit," *Wikipedia*, Feb. 2022. [Online]. Available: https://en.wikipedia.org/w/index.php?title=Application-specific_integrated_circuit&oldid=1074023806 (visited on 04/07/2022).
- [23] "Leela Chess Zero," *Wikipedia*, Apr. 2022. [Online]. Available: https://en.wikipedia.org/w/index.php?title=Leela_Chess_Zero&oldid=1080962634 (visited on 04/08/2022).
- [24] "Leela Zero," *Wikipedia*, Oct. 2021. [Online]. Available: https://en.wikipedia.org/w/index.php?title=Leela_Zero&oldid=1049955001 (visited on 04/08/2022).

- [25] *Lc0*, LCZero, Apr. 2022. [Online]. Available: <https://github.com/LeelaChessZero/lc0> (visited on 04/08/2022).
- [26] *Python-chess: A chess library for Python — python-chess 1.9.0 documentation*. [Online]. Available: <https://python-chess.readthedocs.io/en/latest/> (visited on 04/08/2022).
- [27] “Forsyth–Edwards Notation,” *Wikipedia*, Feb. 2022. [Online]. Available: https://en.wikipedia.org/w/index.php?title=Forsyth%E2%80%93Edwards_Notation&oldid=1069540048 (visited on 04/09/2022).
- [28] *Graphviz*. [Online]. Available: <https://graphviz.org/> (visited on 04/09/2022).
- [29] S. Bodenstein, *AlphaZero*, Sep. 2019. [Online]. Available: <https://sebastianbodenstein.net/post/alphazero/> (visited on 04/09/2022).
- [30] *Lichess.org open database*. [Online]. Available: <https://database.lichess.org/> (visited on 04/09/2022).
- [31] A. Adefokun, *Chess-board ahira-justice*, Feb. 2022. [Online]. Available: <https://github.com/ahira-justice/chess-board> (visited on 04/10/2022).
- [32] *PyGame*. [Online]. Available: <https://www.pygame.org/news> (visited on 04/10/2022).
- [33] zjeffer, *Chess-board zjeffer*, Jan. 2022. [Online]. Available: <https://github.com/zjeffer/chess-board> (visited on 04/10/2022).
- [34] A. Adefokun, *Chess-board pul request*, Feb. 2022. [Online]. Available: <https://github.com/ahira-justice/chess-board/pull/5> (visited on 04/10/2022).
- [35] zjeffer, *Chess engine with Deep Reinforcement learning*, Feb. 2022. [Online]. Available: <https://github.com/zjeffer/chess-deep-rl> (visited on 04/09/2022).
- [36] “First-move advantage in chess,” *Wikipedia*, Mar. 2022. [Online]. Available: https://en.wikipedia.org/w/index.php?title=First-move_advantage_in_chess&oldid=1080096428 (visited on 04/10/2022).
- [37] zjeffer, *Chess-deep-rl-cpp*, Apr. 2022. [Online]. Available: <https://github.com/zjeffer/chess-deep-rl-cpp> (visited on 04/10/2022).
- [38] Thomas Moerland – *Postdoc, Leiden University, The Netherlands*. [Online]. Available: <https://thomasmoerland.nl/> (visited on 04/13/2022).
- [39] “Top Chess Engine Championship,” *Wikipedia*, Mar. 2022. [Online]. Available: https://en.wikipedia.org/w/index.php?title=Top_Chess_Engine_Championship&oldid=1077985240 (visited on 04/12/2022).
- [40] S. R. |. 2. December and 2018, *Has Google cracked the data center cooling problem with AI?* May 2020. [Online]. Available: <https://techwireasia.com/2020/05/has-google-cracked-the-data-centre-cooling-problem-with-ai/> (visited on 04/13/2022).
- [41] *How AI helps better manage and run data centers*, Dec. 2018. [Online]. Available: <https://techwireasia.com/2018/12/how-ai-helps-better-manage-and-run-data-centers/> (visited on 04/14/2022).
- [42] “Reinforcement learning,” *Wikipedia*, Apr. 2022. [Online]. Available: https://en.wikipedia.org/w/index.php?title=Reinforcement_learning&oldid=1081715019 (visited on 04/14/2022).
- [43] nbro, *Answer to "Are there any applications of reinforcement learning other than games?"* Nov. 2020. [Online]. Available: <https://ai.stackexchange.com/a/24355/54037> (visited on 04/13/2022).
- [44] N. T. Blog, *Artwork Personalization at Netflix*, Dec. 2017. [Online]. Available: <https://netflixtechblog.com/artwork-personalization-c589f074ad76> (visited on 04/13/2022).
- [45] P. Surmenok, *Contextual Bandits and Reinforcement Learning*, Oct. 2017. [Online]. Available: <https://towardsdatascience.com/contextual-bandits-and-reinforcement-learning-6bdfeaece72a> (visited on 04/13/2022).
- [46] *Cloud Computing Services | Google Cloud*. [Online]. Available: <https://cloud.google.com/> (visited on 04/13/2022).
- [47] *Introduction to gRPC*. [Online]. Available: <https://grpc.io/docs/what-is-grpc/introduction/> (visited on 04/15/2022).
- [48] *DeepMind | About*. [Online]. Available: <https://www.deepmind.com/about> (visited on 04/13/2022).
- [49] *AlphaZero: Shedding new light on chess, shogi, and Go*. [Online]. Available: <https://www.deepmind.com/blog/alphazero-shedding-new-light-on-chess-shogi-and-go> (visited on 04/13/2022).
- [50] DukeZhou, *Answer to "Is AlphaZero an example of an AGI?"* Nov. 2018. [Online]. Available: <https://ai.stackexchange.com/a/9170/54037> (visited on 04/13/2022).
- [51] *Email exchange between Tuur Vanhoutte and Dr. Thomas Moerland*, Apr. 22.

- [52] “Deep reinforcement learning,” *Wikipedia*, Mar. 2022. [Online]. Available: https://en.wikipedia.org/w/index.php?title=Deep_reinforcement_learning&oldid=1078792888 (visited on 04/14/2022).

8 Appendix